

Demo: traffic lights recognize

Author: Jim Xie

Date: 2020-08-07

```
In [58]: #-*- coding:utf-8 -*-
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.models import model_from_json
from CNNData import CCNNData,CCNNDataBase,g_input_shape,preprocess
import cv2
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
import keras.backend as K
matplotlib inline
from keras import optimizers
```

Data pre-process

```
In [59]: img1=cv2.imread('./samples/train/1/1-center_2018_08_22_09_29_24_368.jpg')
img2=cv2.imread('./samples/train/2/3-center_2018_08_23_10_41_17_054.jpg')
img3=cv2.imread('./samples/train/3/4-center_2018_08_22_15_06_12_233.jpg')
img4=cv2.imread('./samples/train/4/5-center_2018_08_23_10_41_43_916.jpg')
img5=cv2.imread('./samples/train/5/6-center_2018_08_21_18_23_11_991.jpg')
img6=cv2.imread('./samples/train/6/8-center_2018_08_21_16_19_502.jpg')
img0=cv2.imread('./samples/train/0/1-center_2018_08_22_14_54_39_335.jpg')

img1=cv2.cvtColor(img1, cv2.COLOR_BGR2RGB)
img2=cv2.cvtColor(img2, cv2.COLOR_BGR2RGB)
img3=cv2.cvtColor(img3, cv2.COLOR_BGR2RGB)
img4=cv2.cvtColor(img4, cv2.COLOR_BGR2RGB)
img5=cv2.cvtColor(img5, cv2.COLOR_BGR2RGB)
img6=cv2.cvtColor(img6, cv2.COLOR_BGR2RGB)
img0=cv2.cvtColor(img0, cv2.COLOR_BGR2RGB)

fig,ax=plt.subplots(3,3,figsize=(60,40))
plt.subplot(3,3,1)
plt.imshow(img1)
plt.subplot(3,3,2)
plt.imshow(img2)
plt.subplot(3,3,3)
plt.imshow(img3)
plt.subplot(3,3,4)
plt.imshow(img4)
plt.subplot(3,3,5)
plt.imshow(img5)
plt.subplot(3,3,6)
plt.imshow(img6)
plt.subplot(3,3,7)
plt.imshow(img0)
```

```
Out[59]: <matplotlib.image.AxesImage at 0x7f961f954550>
```

```
In [60]: img1_1=preprocess(img1)
img2_1=preprocess(img2)
img3_1=preprocess(img3)
img4_1=preprocess(img4)
img5_1=preprocess(img5)
img6_1=preprocess(img6)
fig,ax=plt.subplots(2,3,figsize=(10,6))
plt.subplot(2,3,1)
plt.imshow(img1_1,cmap='gray')
plt.subplot(2,3,2)
plt.imshow(img2_1,cmap='gray')
plt.subplot(2,3,3)
plt.imshow(img3_1,cmap='gray')
plt.subplot(2,3,4)
plt.imshow(img4_1,cmap='gray')
plt.subplot(2,3,5)
plt.imshow(img5_1,cmap='gray')
plt.subplot(2,3,6)
plt.imshow(img6_1,cmap='gray')
```

```
Out[60]: <matplotlib.image.AxesImage at 0x7f960b184e0>
```

Load dataset

```
In [61]: g_test_train=CCNNData()
g_test_train.LoadTrain()
g_test_data=CCNNData()
g_test_data.LoadVerify()
print("Load finished")

adding train sample tag = .DS_Store folder = ./samples/Train/.DS_Store/
adding train sample tag = 0 folder = ./samples/Train/0/
adding train sample tag = 6 folder = ./samples/Train/6/
adding train sample tag = 1 folder = ./samples/Train/1/
adding train sample tag = 4 folder = ./samples/Train/4/
adding train sample tag = 3 folder = ./samples/Train/3/
adding train sample tag = 2 folder = ./samples/Train/2/
adding train sample tag = 5 folder = ./samples/Train/5/
adding verify sample tag = .DS_Store folder = ./samples/verify/.DS_Store/
adding verify sample tag = 0 folder = ./samples/verify/0/
adding verify sample tag = 6 folder = ./samples/verify/6/
adding verify sample tag = 1 folder = ./samples/verify/1/
adding verify sample tag = 4 folder = ./samples/verify/4/
adding verify sample tag = 3 folder = ./samples/verify/3/
adding verify sample tag = 2 folder = ./samples/verify/2/
adding verify sample tag = 5 folder = ./samples/verify/5/
Load finished
```

Begin Train

```
In [76]: class CCNNTrain:
def __init__(self):
    self.m_model=Sequential()
    self.m_learn_rate=0.0005 #0.5
    self.m_adam=optimizers.Adam(lr=self.m_learn_rate)

def CreateModel(self,n_classes):
    self.m_model.add(Conv2D(28,(3,3),activation='sigmoid',input_shape=g_input_shape))
    self.m_model.add(Conv2D(28,(3,3),activation='sigmoid'))
    self.m_model.add(MaxPooling2D(pool_size=(2,2)))
    self.m_model.add(Dropout(0.25))
    self.m_model.add(Flatten())
    self.m_model.add(Dense(128,activation='sigmoid'))
    self.m_model.add(Dropout(0.5))
    self.m_model.add(Dense(n_classes,activation='softmax'))
    self.m_model.compile(loss='categorical_crossentropy',optimizer=self.m_adam,metrics=['accuracy'])

def Fit(self,x_train,y_train):
    return self.m_model.fit(x_train,y_train,batch_size=32,epochs=50,verbose=1)

def Evaluate(self,x_test,y_test):
    score=self.m_model.evaluate(x_test,y_test,verbose=1)
    return score

def Predict(self,image):
    x=CCNNDataBase().GetPredict(image)
    if x is None:
        return 0
    predictions=self.m_model.predict(x)
    for pred in predictions.tolist():
        confidence=max(pred)
        if confidence<0.5:
            return 0
        return pred.index(max(pred))

def Verify(self,data):
    predictions=self.m_model.predict(data)
    for pred in predictions.tolist():
        confidence=max(pred)
        if confidence<0.5:
            return 0
        return pred.index(max(pred))

def Load(self):
    with open('./model/model.json','r') as json_file:
        loaded_model_json=json_file.read()
    self.m_model=model_from_json(loaded_model_json)
    self.m_model.load_weights('./model/model.h5')
    self.m_model.compile(loss='categorical_crossentropy',optimizer=self.m_adam,metrics=['accuracy'])
    return self.m_model

def Save(self):
    model_json=self.m_model.to_json()
    with open('./model/model.json','w') as json_file:
        json_file.write(model_json)
    #serialize weights to HDF5
    self.m_model.save_weights('./model/model.h5')
```

```
In [77]: n_train_class,y_train_data,x_train_data=g_test_train.GetTrain()
cnn=CCNNTrain()
cnn.CreateModel(n_train_class)
print("training")
measure=cnn.Fit(x_train_data,y_train_data)
cnn.Save()

training
Epoch 1/50
278/278 [=====] - 9s 34ms/step - loss: 1.2089 - accuracy: 0.5425
Epoch 2/50
278/278 [=====] - 9s 34ms/step - loss: 0.6283 - accuracy: 0.7681
Epoch 3/50
278/278 [=====] - 9s 33ms/step - loss: 0.5715 - accuracy: 0.7833
Epoch 4/50
278/278 [=====] - 10s 34ms/step - loss: 0.5415 - accuracy: 0.7906
Epoch 5/50
278/278 [=====] - 10s 35ms/step - loss: 0.5221 - accuracy: 0.8003
Epoch 6/50
278/278 [=====] - 9s 34ms/step - loss: 0.5026 - accuracy: 0.8109
Epoch 7/50
278/278 [=====] - 9s 34ms/step - loss: 0.4868 - accuracy: 0.8139
Epoch 8/50
278/278 [=====] - 9s 34ms/step - loss: 0.4719 - accuracy: 0.8234
Epoch 9/50
278/278 [=====] - 10s 34ms/step - loss: 0.4597 - accuracy: 0.8243
Epoch 10/50
278/278 [=====] - 9s 33ms/step - loss: 0.4495 - accuracy: 0.8309
Epoch 11/50
278/278 [=====] - 9s 34ms/step - loss: 0.4355 - accuracy: 0.8335
Epoch 12/50
278/278 [=====] - 9s 33ms/step - loss: 0.4243 - accuracy: 0.8404
Epoch 13/50
278/278 [=====] - 10s 35ms/step - loss: 0.4152 - accuracy: 0.8408
Epoch 14/50
278/278 [=====] - 9s 33ms/step - loss: 0.4040 - accuracy: 0.8440
Epoch 15/50
278/278 [=====] - 10s 34ms/step - loss: 0.3933 - accuracy: 0.8500
Epoch 16/50
278/278 [=====] - 9s 31ms/step - loss: 0.3852 - accuracy: 0.8556
Epoch 17/50
278/278 [=====] - 8s 30ms/step - loss: 0.3701 - accuracy: 0.8616
Epoch 18/50
278/278 [=====] - 10s 36ms/step - loss: 0.3650 - accuracy: 0.8626
Epoch 19/50
278/278 [=====] - 10s 37ms/step - loss: 0.3549 - accuracy: 0.8682
Epoch 20/50
278/278 [=====] - 10s 37ms/step - loss: 0.3475 - accuracy: 0.8701
Epoch 21/50
278/278 [=====] - 10s 37ms/step - loss: 0.3365 - accuracy: 0.8763
Epoch 22/50
278/278 [=====] - 10s 36ms/step - loss: 0.3285 - accuracy: 0.8775
Epoch 23/50
278/278 [=====] - 9s 33ms/step - loss: 0.3285 - accuracy: 0.8770
Epoch 24/50
278/278 [=====] - 9s 32ms/step - loss: 0.3083 - accuracy: 0.8842
Epoch 25/50
278/278 [=====] - 9s 33ms/step - loss: 0.3038 - accuracy: 0.8844
Epoch 26/50
278/278 [=====] - 10s 35ms/step - loss: 0.2952 - accuracy: 0.8908
Epoch 27/50
278/278 [=====] - 10s 33ms/step - loss: 0.2825 - accuracy: 0.8949
Epoch 28/50
278/278 [=====] - 10s 35ms/step - loss: 0.2806 - accuracy: 0.8952
Epoch 29/50
278/278 [=====] - 10s 35ms/step - loss: 0.2687 - accuracy: 0.8999
Epoch 30/50
278/278 [=====] - 10s 35ms/step - loss: 0.2560 - accuracy: 0.9066
Epoch 31/50
278/278 [=====] - 10s 34ms/step - loss: 0.2610 - accuracy: 0.9037
Epoch 32/50
278/278 [=====] - 10s 34ms/step - loss: 0.2481 - accuracy: 0.9099
Epoch 33/50
278/278 [=====] - 10s 35ms/step - loss: 0.2409 - accuracy: 0.9109
Epoch 34/50
278/278 [=====] - 10s 34ms/step - loss: 0.2402 - accuracy: 0.9151
Epoch 35/50
278/278 [=====] - 9s 34ms/step - loss: 0.2292 - accuracy: 0.9169
Epoch 36/50
278/278 [=====] - 9s 34ms/step - loss: 0.2106 - accuracy: 0.9242
Epoch 37/50
278/278 [=====] - 10s 35ms/step - loss: 0.2131 - accuracy: 0.9242
Epoch 38/50
278/278 [=====] - 10s 34ms/step - loss: 0.2014 - accuracy: 0.9318
Epoch 39/50
278/278 [=====] - 10s 35ms/step - loss: 0.2021 - accuracy: 0.9274
Epoch 40/50
278/278 [=====] - 9s 34ms/step - loss: 0.1965 - accuracy: 0.9313
Epoch 41/50
278/278 [=====] - 10s 35ms/step - loss: 0.1927 - accuracy: 0.9335
Epoch 42/50
278/278 [=====] - 9s 33ms/step - loss: 0.1823 - accuracy: 0.9349
Epoch 43/50
278/278 [=====] - 10s 35ms/step - loss: 0.1762 - accuracy: 0.9406
Epoch 44/50
278/278 [=====] - 10s 37ms/step - loss: 0.1727 - accuracy: 0.9403
Epoch 45/50
278/278 [=====] - 10s 37ms/step - loss: 0.1752 - accuracy: 0.9389
Epoch 46/50
278/278 [=====] - 9s 34ms/step - loss: 0.1670 - accuracy: 0.9373
Epoch 47/50
278/278 [=====] - 10s 37ms/step - loss: 0.1581 - accuracy: 0.9438
Epoch 48/50
278/278 [=====] - 10s 37ms/step - loss: 0.1485 - accuracy: 0.9480
Epoch 49/50
278/278 [=====] - 10s 37ms/step - loss: 0.1469 - accuracy: 0.9496
Epoch 50/50
278/278 [=====] - 10s 36ms/step - loss: 0.1558 - accuracy: 0.9474
```

```
In [78]: plt.style.use({'figure.figsize':(12,6)})
df=pd.DataFrame()
df['loss']=measure.history['loss']
df['accuracy']=measure.history['accuracy']
#df['val_loss']=measure.history['val_loss']
#df['val_accuracy']=measure.history['val_accuracy']
xticks=range(len(df['accuracy']))
df.plot(xticks=xticks)
plt.ylabel('loss/accuracy')
plt.legend()
plt.show()
```

Verify

```
In [79]: cnn=CCNNTrain()
cnn.Load()
count=g_test_data.m_verify.GetSize()
NotFound=[]
x_list=[]
label_list=[]
pred_list=[]
result={}
for i in range(count):
    x=g_test_data.m_verify.GetVerifyData(i)
    y=g_test_data.m_verify.GetVerifyResult(i)
    f=g_test_data.m_verify.GetVerifyFile(i)
    v=cnn.Verify(x)
    if str(v)=='0':
        NotFound.append(f)
        continue
    if not v in result:
        result[v]={}
    img=cv2.imread(f)
    img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
    img=preprocess(img)
    if len(img)>0:
        result[v].append(img)
        label_list.append(int(y))
        pred_list.append(int(v))

label_list=np.array(label_list)
pred_list=np.array(pred_list)
```

```
In [80]: true_label=label_list
predict_label=pred_list
pd.crosstab(true_label,predict_label,rownames=['label'],colnames=['predict'])
```

```
Out[80]:
```

	label	1	2	3	4	5	6
label	0	0	2	0	0	0	0
1	83	0	2	0	2	5	
2	1	84	1	0	0	5	
3	0	3	57	1	1	8	
4	0	0	85	1	0		
5	2	0	0	47	1		
6	4	13	1	1	2	43	

```
In [87]: from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score
confusion_matrix=metrics.confusion_matrix(true_label,predict_label)
sns.heatmap(confusion_matrix,annot=True,annot_kws={"size":19},cmap="Blues")
```

```
Out[87]: <AxesSubplot:~>
```

```
In [88]: from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score
print('-----Weighted-----')
print('Weighted precision',precision_score(true_label,predict_label,average='weighted'))
print('Weighted recall',recall_score(true_label,predict_label,average='weighted'))
print('Weighted f1-score',f1_score(true_label,predict_label,average='weighted'))
print('-----Macro-----')
print('Macro precision',precision_score(true_label,predict_label,average='macro'))
print('Macro recall',recall_score(true_label,predict_label,average='macro'))
print('Macro f1-score',f1_score(true_label,predict_label,average='macro'))
print('-----Micro-----')
print('Micro precision',precision_score(true_label,predict_label,average='micro'))
print('Micro recall',recall_score(true_label,predict_label,average='micro'))
print('Micro f1-score',f1_score(true_label,predict_label,average='micro'))

-----Weighted-----
Weighted precision: 0.8657817979190768
Weighted recall: 0.8682352941176471
Weighted f1-score: 0.865704368460885
-----Macro-----
Macro precision: 0.7464901405891202
Macro recall: 0.7476506294498533
Macro f1-score: 0.7459136585945193
-----Micro-----
Micro precision: 0.8682352941176471
Micro recall: 0.8682352941176471
Micro f1-score: 0.8682352941176471
```

```
In [89]: def ShowImg(img_list):
rows=3
if len(img_list)<24:
    rows=len(img_list)//8
fig,ax=plt.subplots(rows,8,figsize=(10,5))
for j in range(0,rows):
    for i in range(1,9):
        plt.subplot(rows,8,i+j*8)
        plt.imshow(img_list[i+j*8],cmap='gray')
        plt.axis('off')
```

```
In [84]: ShowImg(result[3])
```

```
In [85]: def ShowLayer(data,num_layer=1):
data=np.expand_dims(data,axis=0)
layer=K.function([cnn.m_model.layers[0].input],[cnn.m_model.layers[num_layer].output])
f1=layer([data])[0]
num=f1.shape[-1]
plt.figure(figsize=(10,10))
for i in range(num):
    plt.subplot(8,8,i+1)
    plt.imshow(f1[0,i,i],cmap='gray')
    plt.axis('off')
plt.show()
```

```
In [86]: ShowLayer(result[2][0],2)
```

```
In [ ]:
```

```
In [ ]:
```