Practice of AI

C2: Machine learning & Data analyze

谢文伟 (Jim Xie)

2020/3/6

Outline

1. Goal
2. ML workflow introduction
3. Basic math knowledge introduction
4. Time series forecasting demo
5. Brief summary

Goal

Getting start for data analyze with ML

Demo #1

Is ML Universal ?

Limitations

- NFL
- Smooth
- Boundary

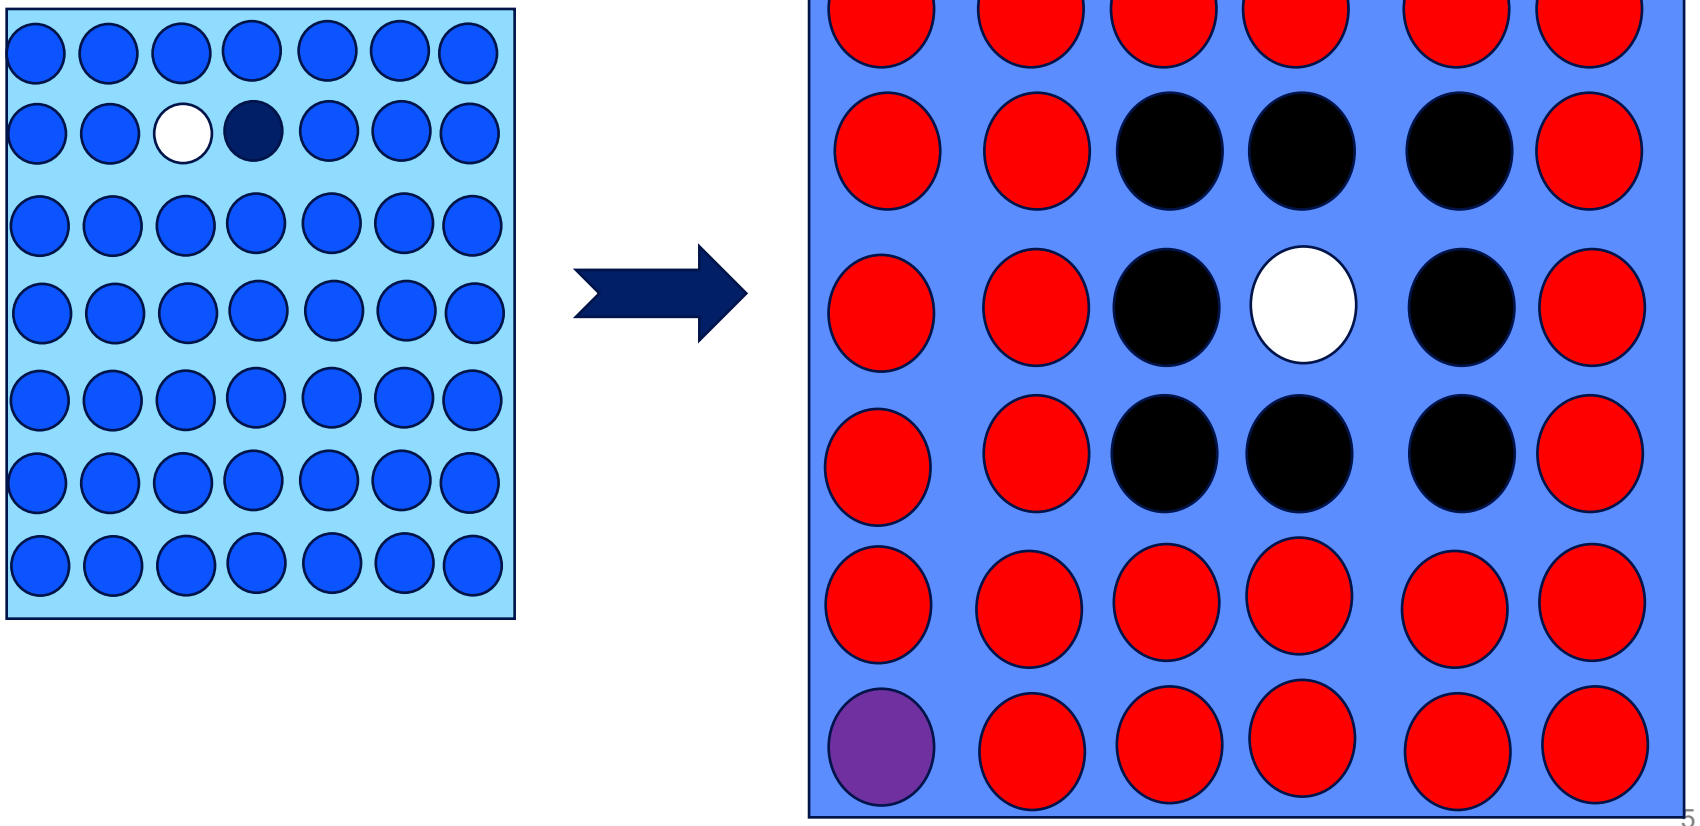

Category by sample

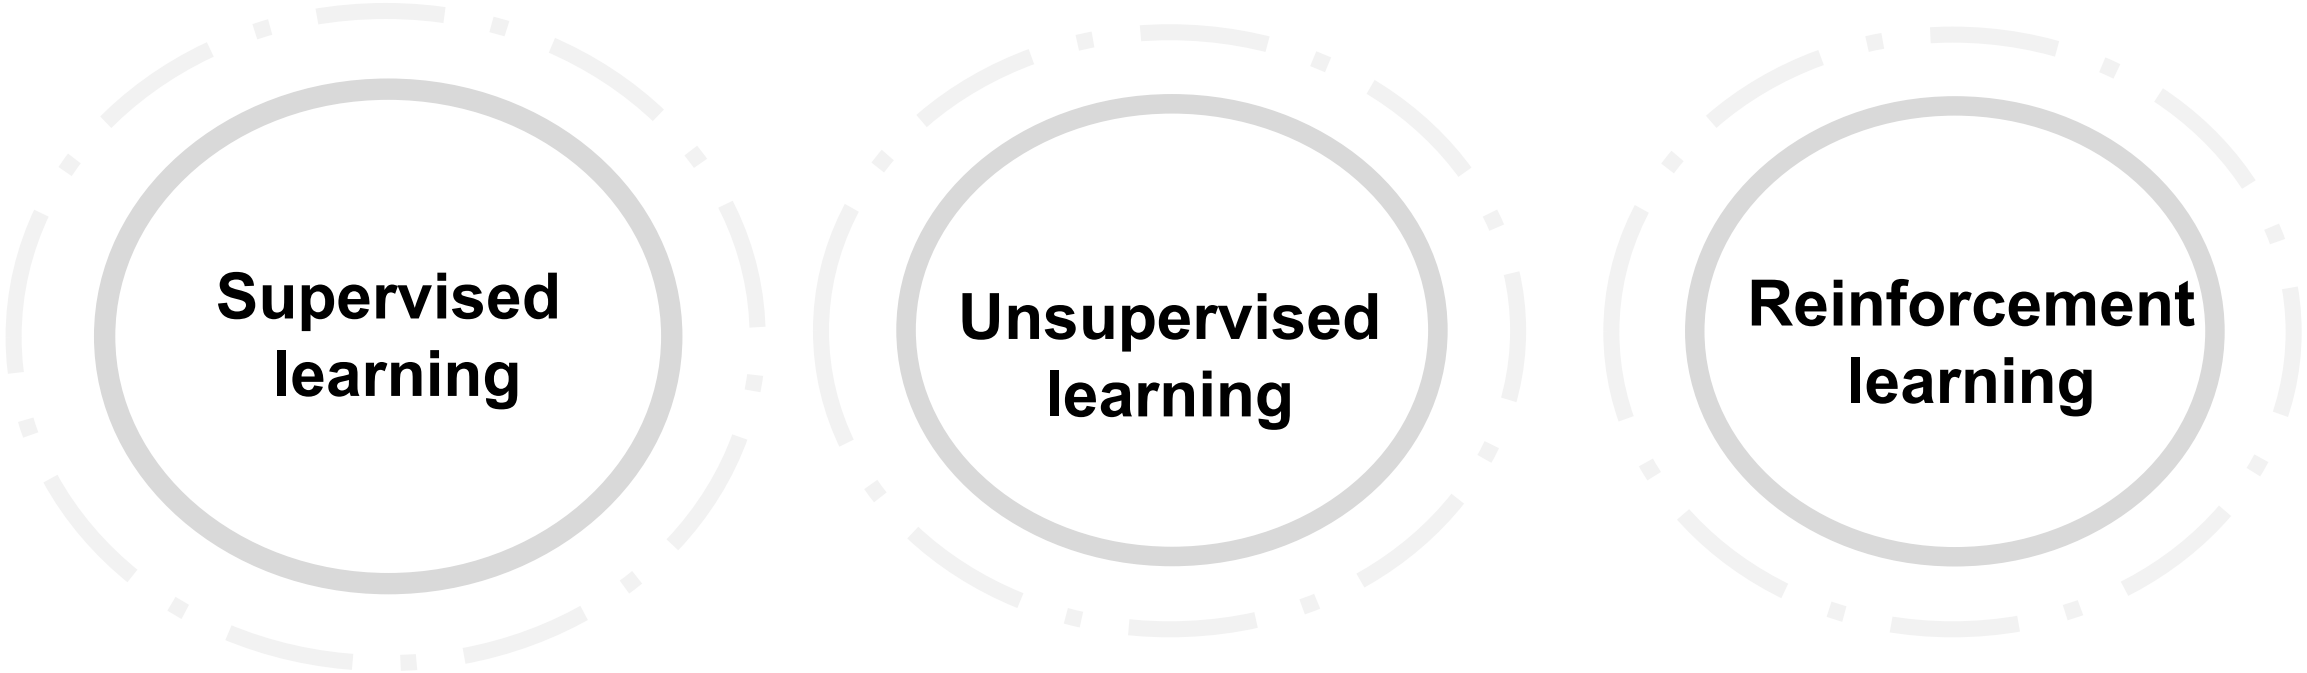

Q: 样本不平衡怎么办？

Purpose

- **Class**

Cat or Dog?

- **Regression**

How much?

Workflow

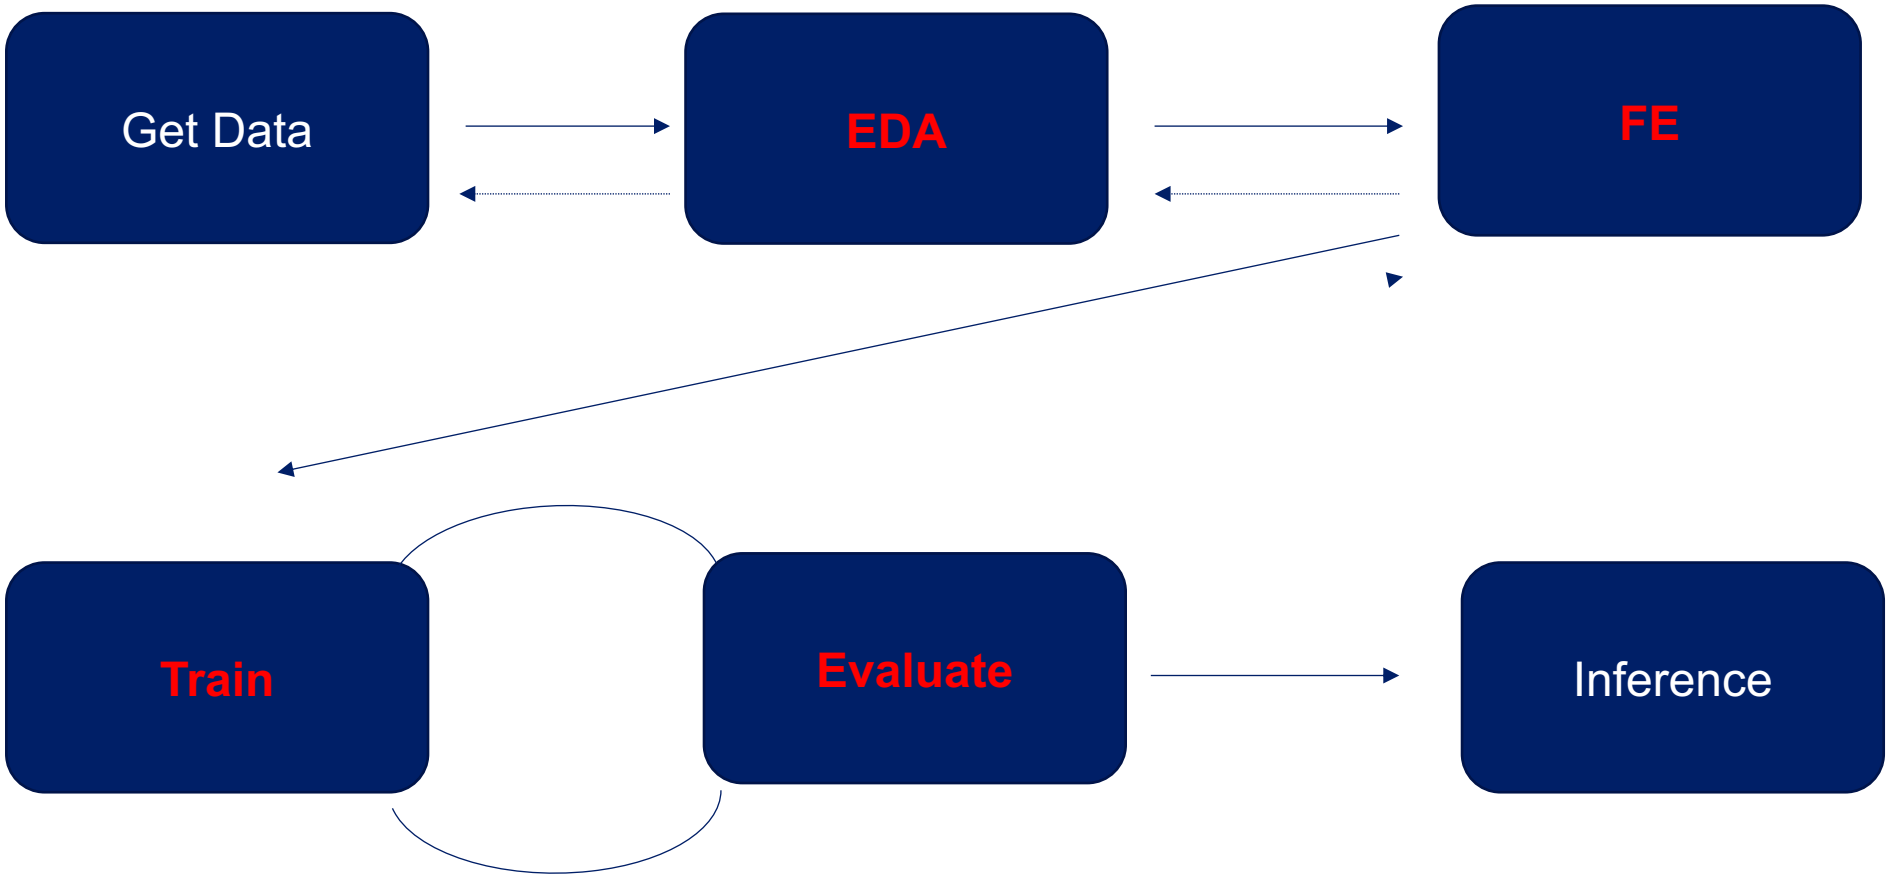

EDA & preprocess

EDA #1

- 检查样本是否合格？
- 样本量有多大？
- 有多少缺失数据？
- ……

- 回归 or 分类？
- 发现潜在的特征
- 数据如何分布？
- ……

EDA #2

- Insight from graph
- Data distribution
- Normalization

Data distribution

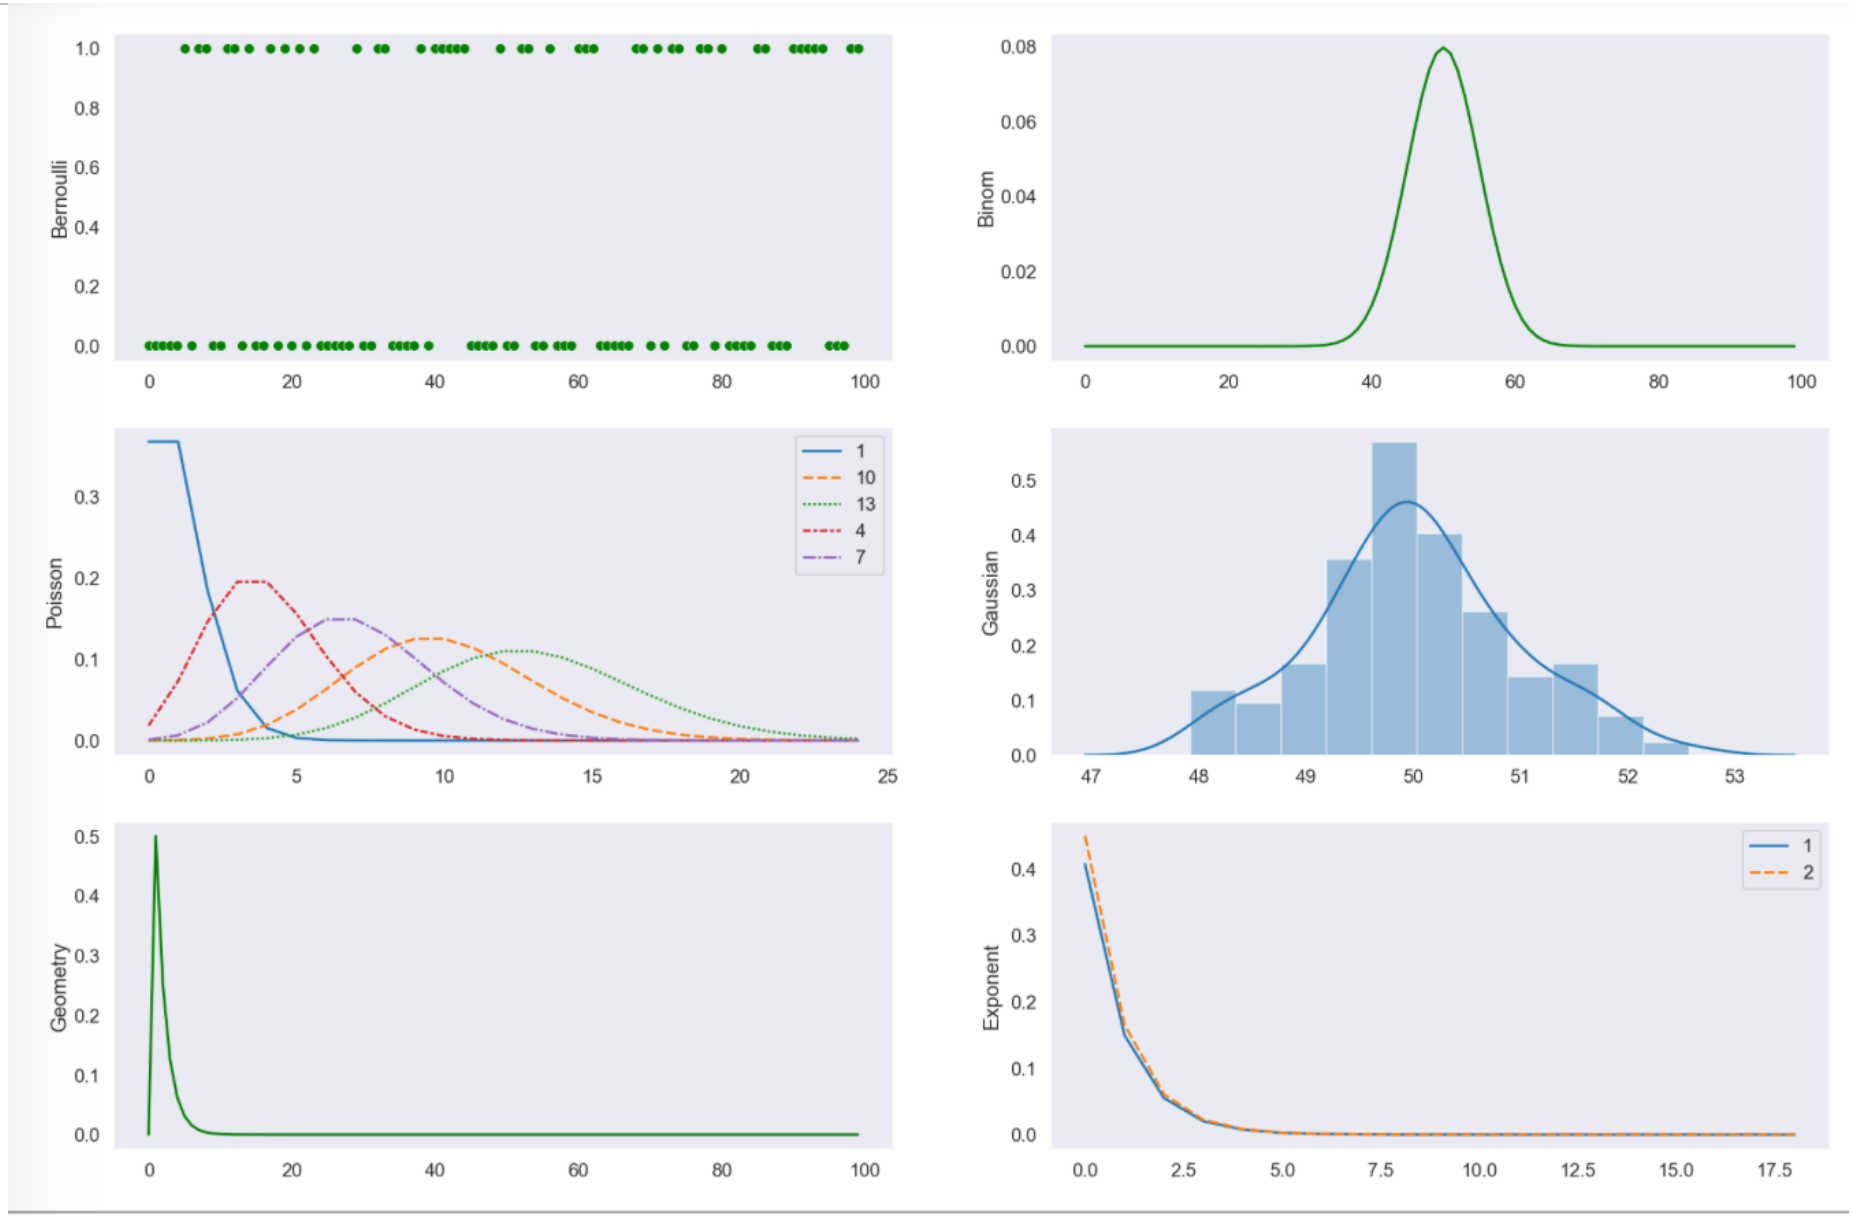

Normalization

- Why ?

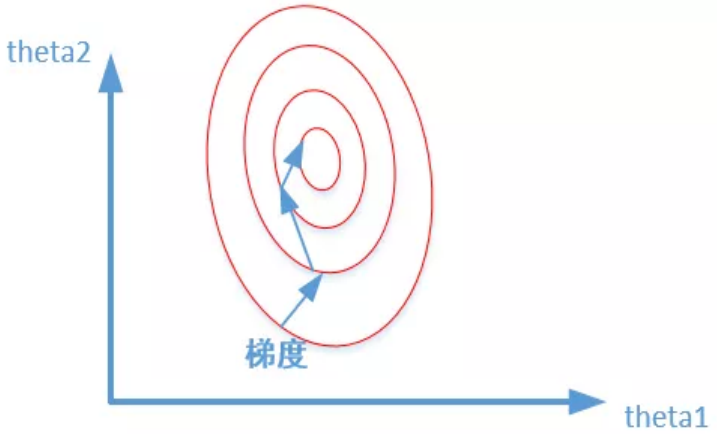

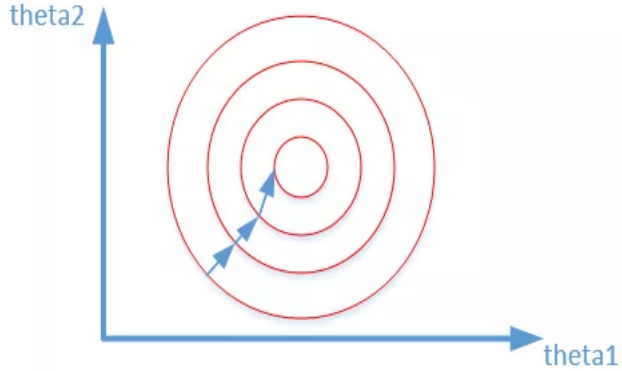

- How ?

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

$$x' = \frac{x - \mu}{\sigma}$$

Feature Engine

FE #1

Why

01

减少噪音

提高模型性能

02

减少维数

降低运算量

03

降低复杂度

增加可解释性

FE #2

- 特征是否发散？
- 特征和目标是否相关？

- Filter Methods
  - Correlation
  - ….
- Embed Methods
  - GA
  - ….
- Wrap methods
  - CNN
  - ….

FE #3

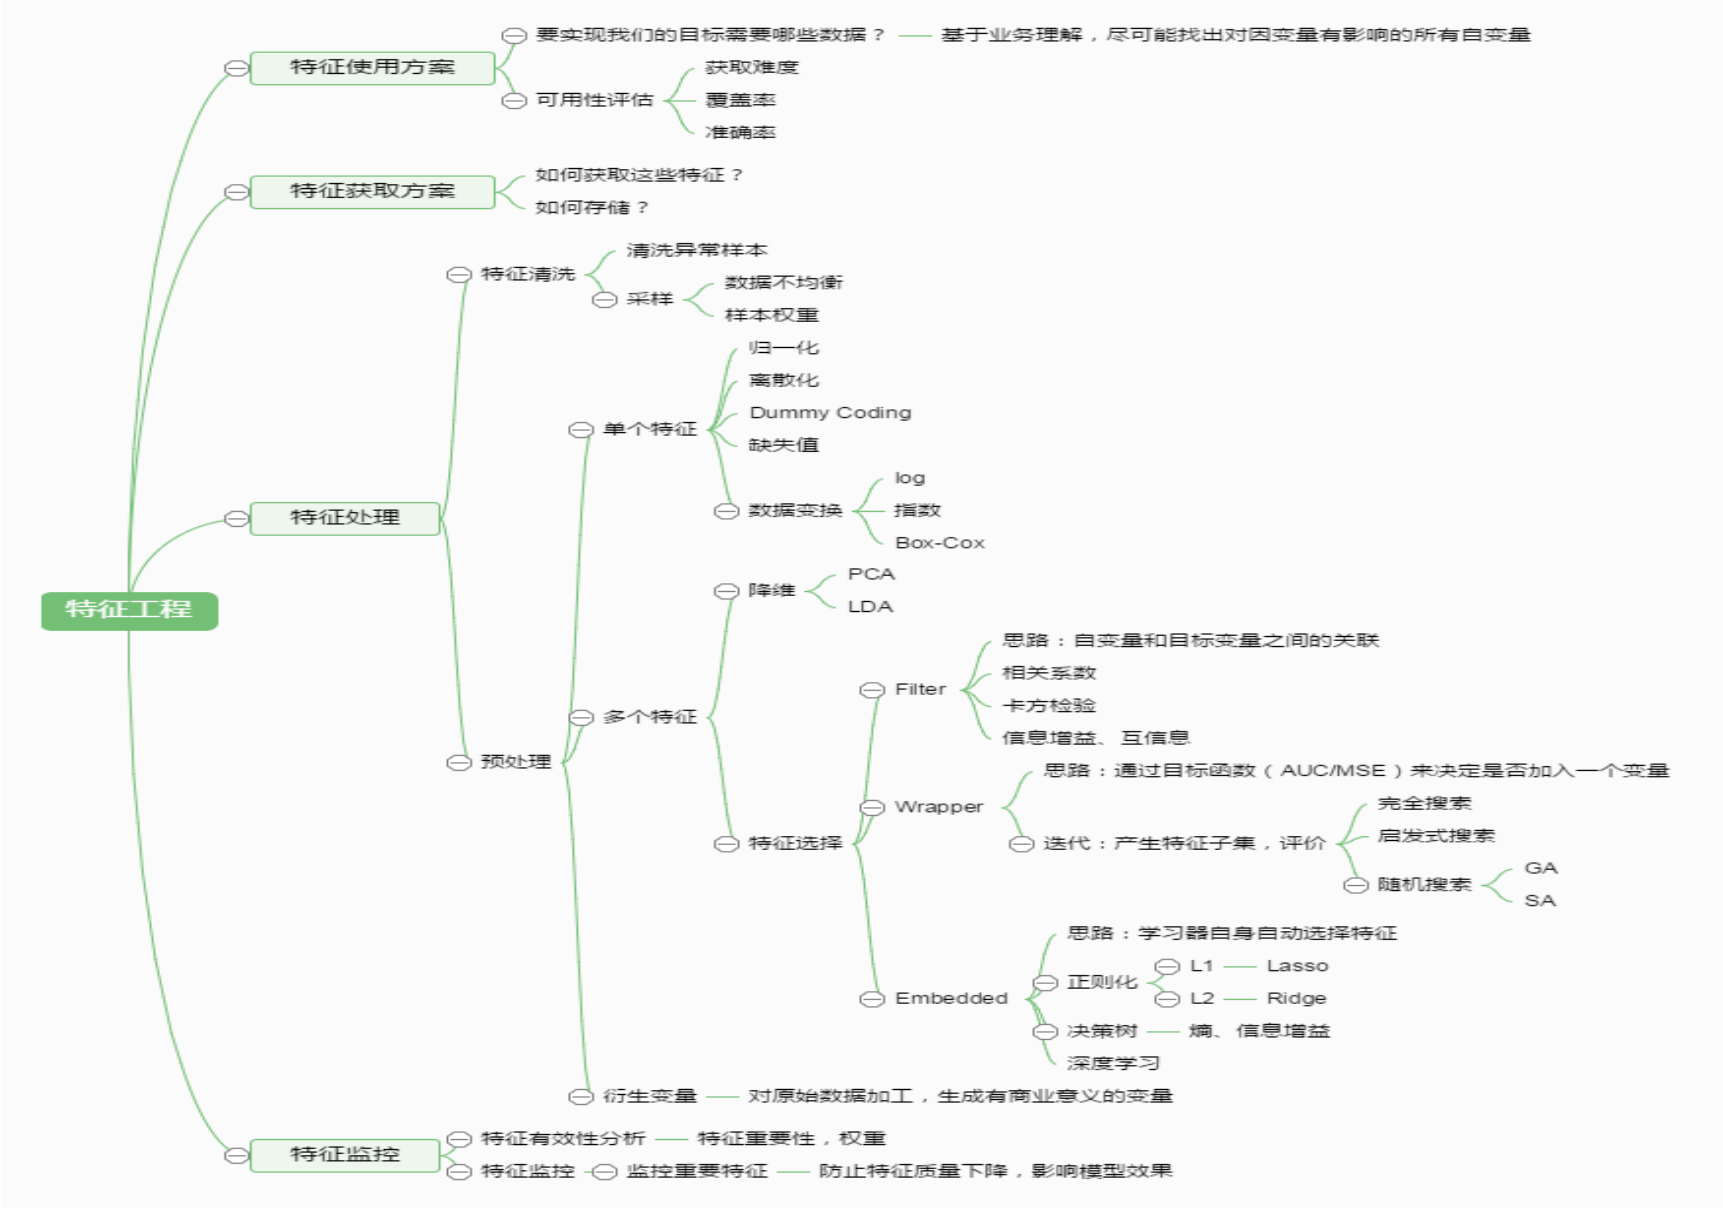
特征工程
特征使用方案
要实现我们的目标需要哪些数据？ —— 基于业务理解，尽可能找出对因变量有影响的所有自变量
可用性评估
获取难度
覆盖率
准确率
特征获取方案
如何获取这些特征？
如何存储？
特征处理
特征清洗
清洗异常样本
采样
数据不均衡
样本权重
预处理
单个特征
归一化
离散化
Dummy Coding
缺失值
数据变换
log
指数
Box-Cox
多个特征
降维
PCA
LDA
特征选择
Filter
思路：自变量和目标变量之间的关联
相关系数
卡方检验
信息增益、互信息
Wrapper
思路：通过目标函数（AUC/MSE）来决定是否加入一个变量
迭代：产生特征子集，评价
完全搜索
启发式搜索
随机搜索
GA
SA
Embedded
思路：学习器自身自动选择特征
正则化
L1 —— Lasso
L2 —— Ridge
决策树 —— 熵、信息增益
深度学习
衍生变量 —— 对原始数据加工，生成有商业意义的变量
特征监控
特征有效性分析 —— 特征重要性，权重
特征监控 —— 监控重要特征 —— 防止特征质量下降，影响模型效果

FE #4

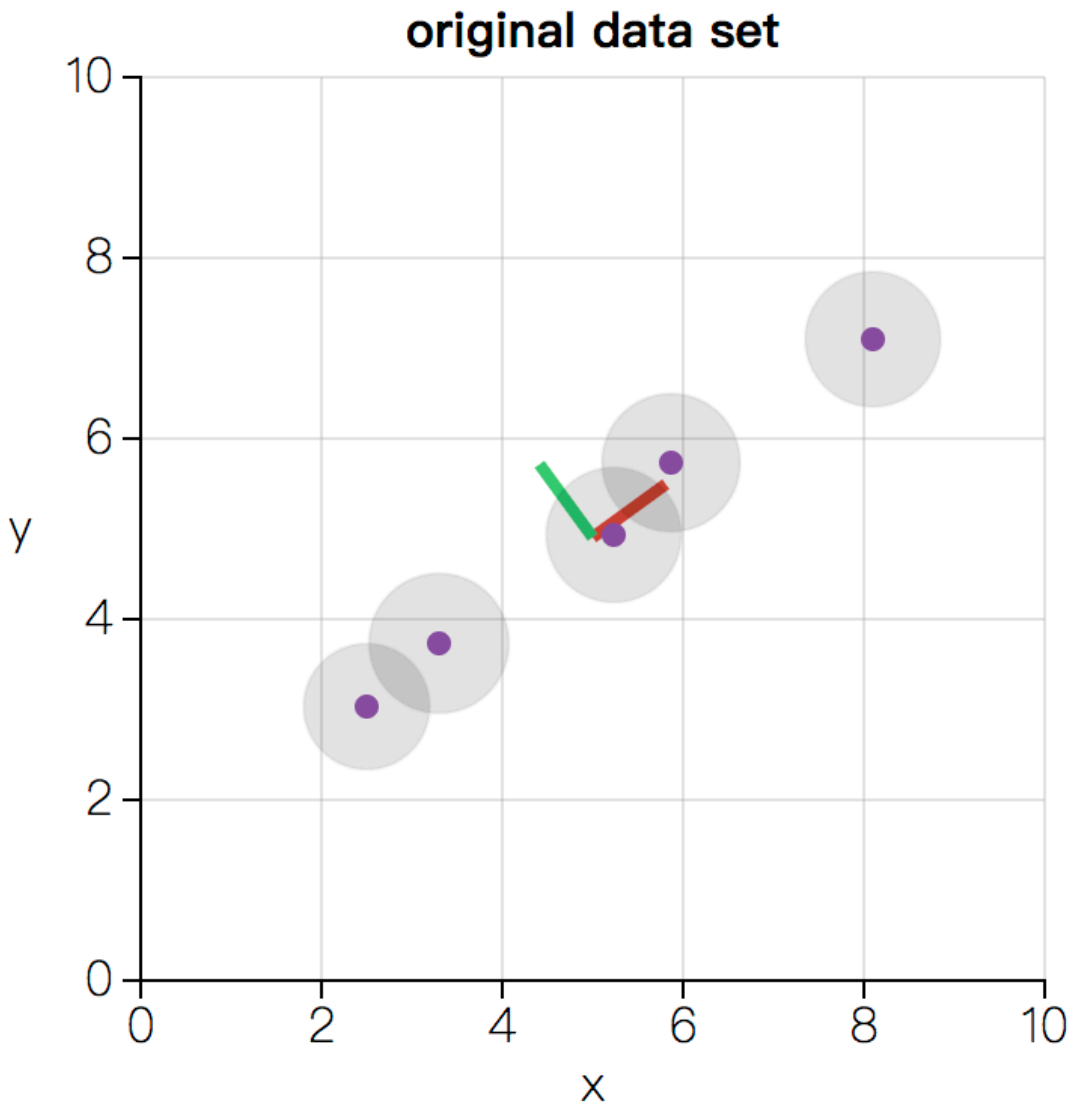
original data set
10
8
6
4
2
0
y
0
2
4
6
8
10
x

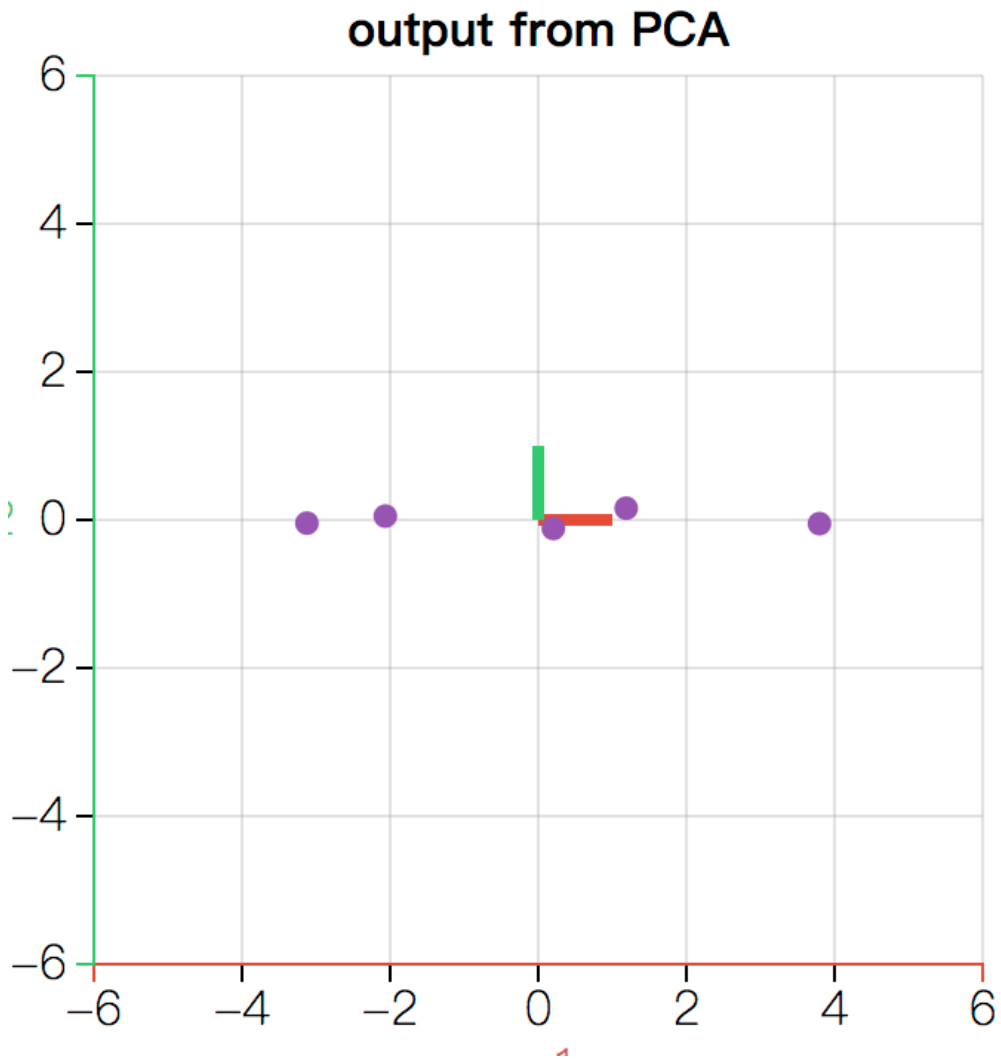
output from PCA
6
4
2
0
−2
−4
−6
−6
−4
−2
0
2
4
6

Train and Evaluation

Model and Evaluation

Models

- **Linear**
- **Neural network**
- DT/SVM/Bayes
- **XNN**/LSTM/GRU
- Boost/XGBoost
- ……..

Evaluate # Regression

Error = | Real – Predict | 

$$RMSE(X,h) = \sqrt{\frac{1}{m}\sum_{i=1}^{m}(h(x_i) - y_i)^2}$$

$$MSE = \frac{1}{m}\sum_{i=1}^{m}(y_i - \hat{y}_i)^2$$

$$MAE(X,h) = \frac{1}{m}\sum_{i=1}^{m}|h(x_i) - y_i|$$

$$SD = \sqrt{\frac{1}{N}\sum_{i=1}^{N}(x_i - avg(x))^2}$$

Evaluate # Class

Confusion Matrix

	Predicted (Positive)	Predicted (Negative)
Actual (Positive)	TP	FN
Actual (Negative)	FP	TN

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Sensitivity} = \text{Recall} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

$$F_1 = 2 \cdot \frac{precision \cdot recall}{precison + recall}$$

ROC

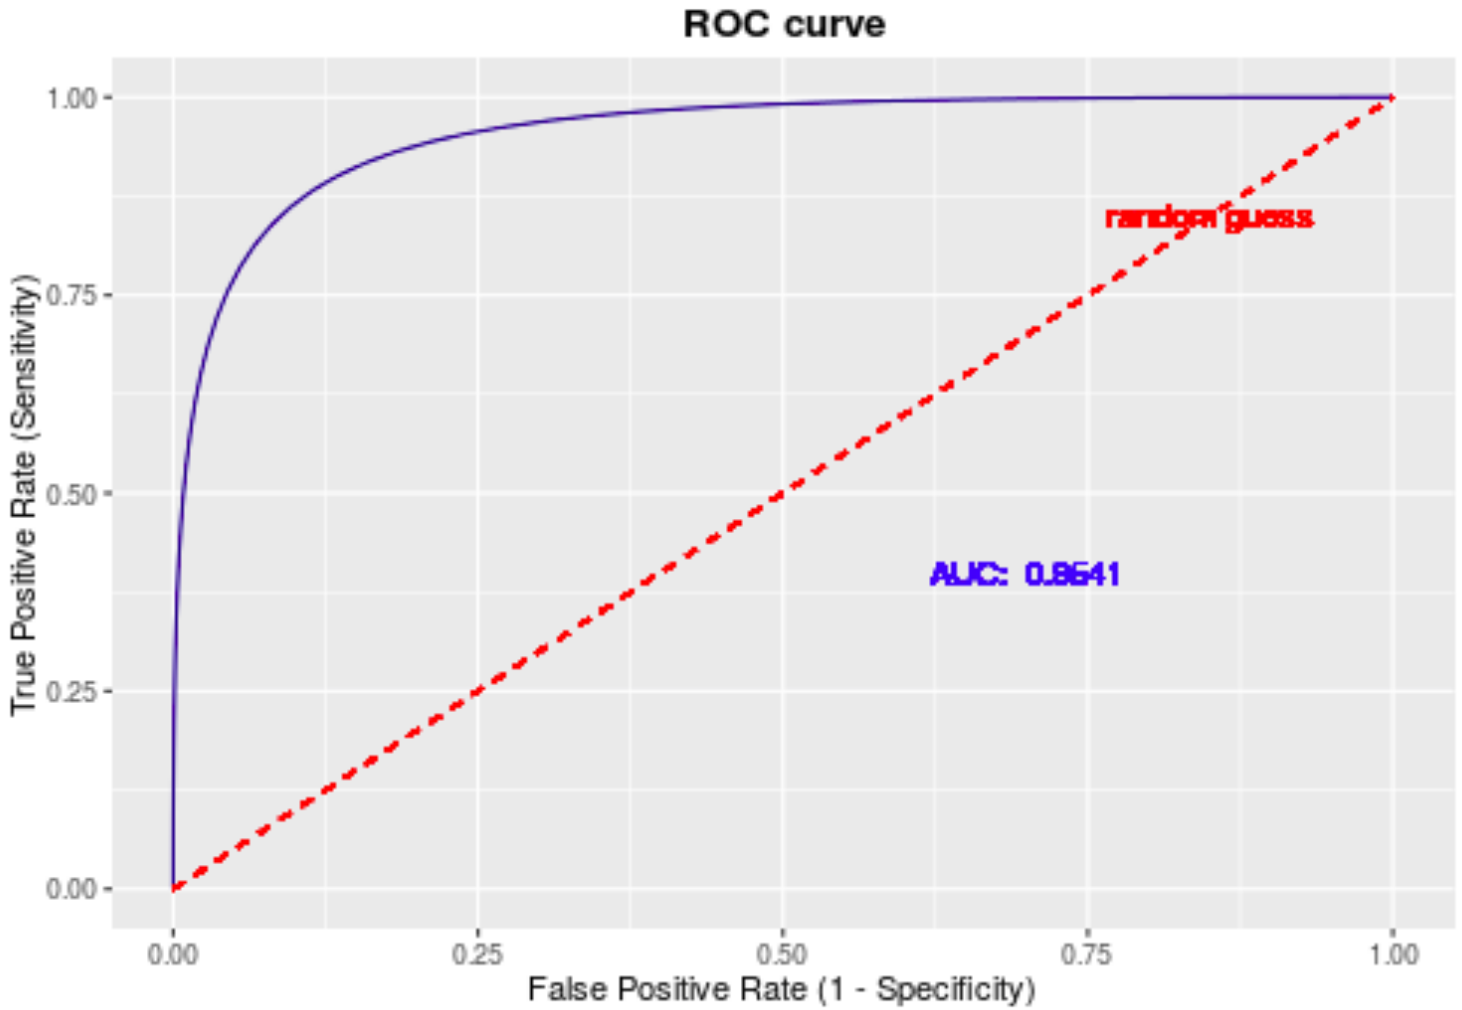
ROC curve
True Positive Rate (Sensitivity)
False Positive Rate (1 - Specificity)
0.00
0.25
0.50
0.75
1.00
random guess
AUC: 0.8641

Underfitting/Overfitting

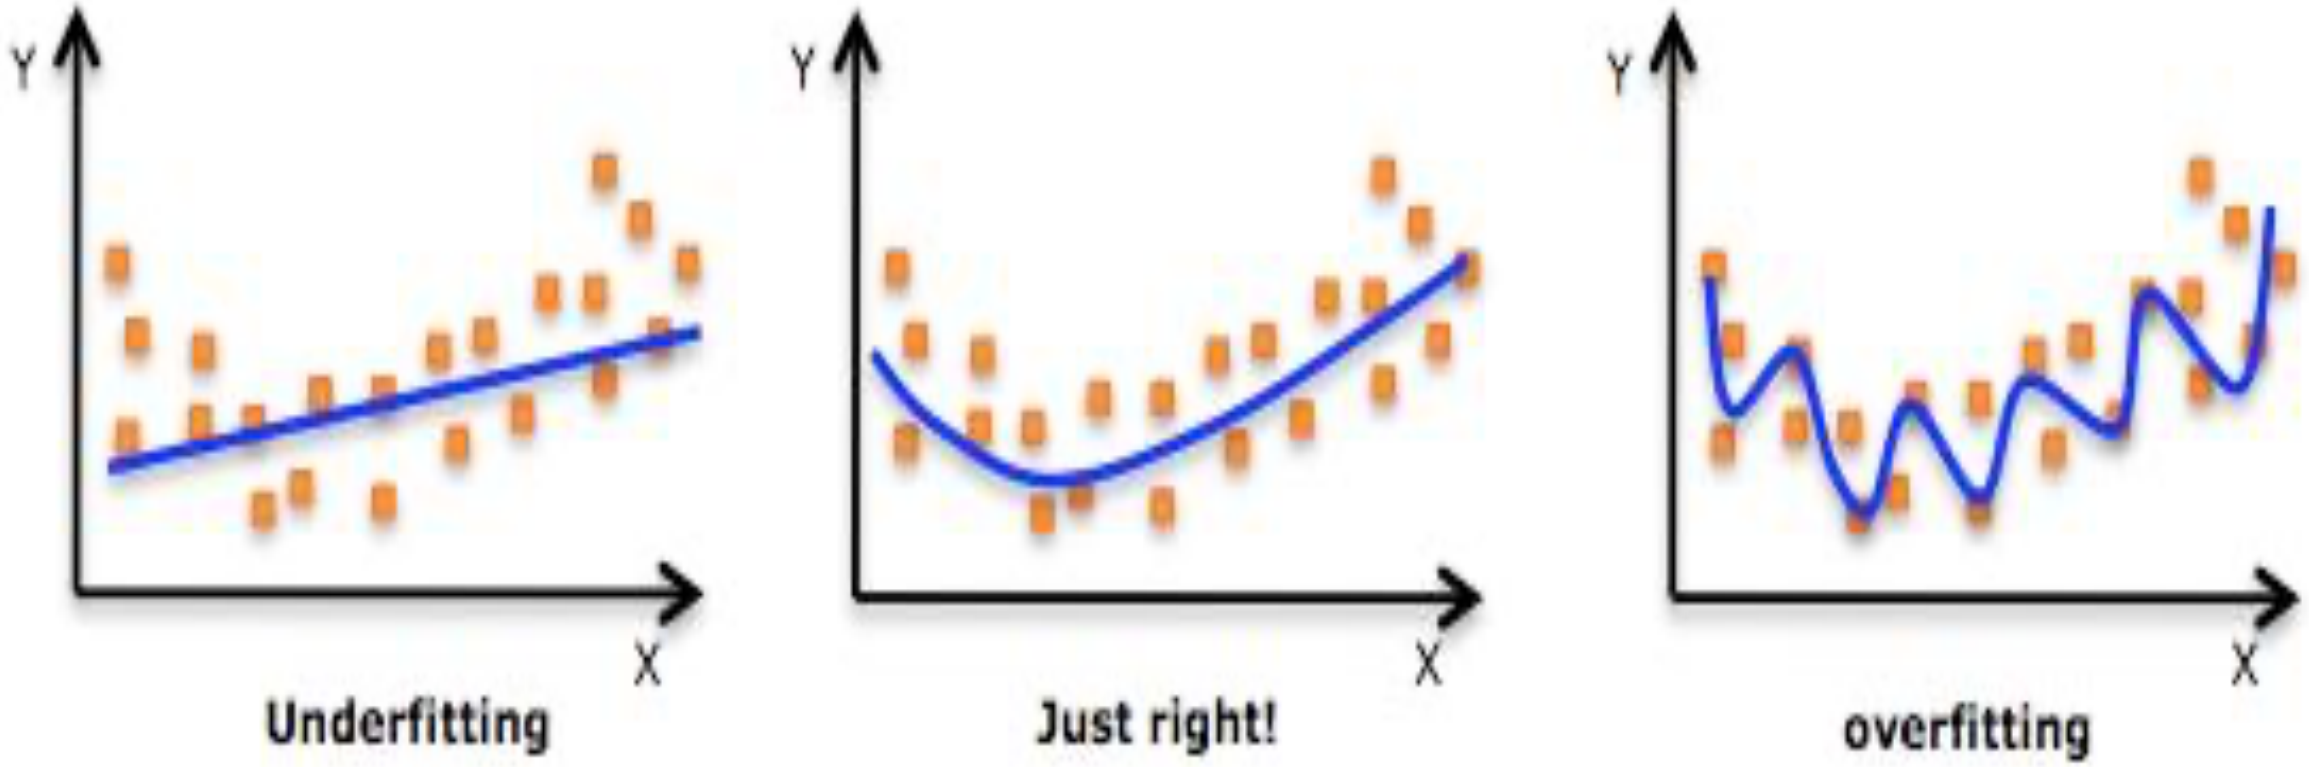

Q: How to do ?

Backlog

Backlog